

KubeVirt on Talos: A Homelab Journey



+



=



Run KubeVirt, they said...

It would be fun, they said...

It took me 7 months.... 😂



Who am i?

Name: Michael Trip

From: Apeldoorn / The Netherlands

- Open Source consultant @ AT Computing
- Kubernetes trainer
- Current consultancy gig: Dutch Tax Administration
- Linux Geek (started in 2004)
- Redhat certs: RHCSA, RHCE and Openshift
- Kubernetes certs: CKAD, CKA and CKS
- Hypervisor experience:
 - VMWare GSX (back in 2008)
 - VMWare ESXi (since 3.5)
 - HyperV
 - Xen
 - Proxmox
 - KubeVirt (obviously)



And.....



Topics of today

- My homelab
- What makes a good hypervisor?
- What is KubeVirt ?
- KubeVirt installation on Talos
- CDI: Containerized Data Importer
- What about.....
 - Vlans and networking?
 - Hyperconverged and shared storage?
 - Live migration?
- Demo time:
 - Creating VMs
 - Live migration
- Takeaways and conclusion



Homelab

- 1 virtual control plane node running Talos 1.7.6 on Proxmox
- 1 Zimaboard running NFS With Debian 12
- 3 worker nodes running on bare metal with Talos 1.7.6
 - 2 HP Elitedesks with 16GB ram, 256GB SSD, single disk
 - 1 HP Prodesk with 16 GB ram, 256GB SSD, single disk
- Endgoal for my project: Have a enterprise-like virtualization cluster

```
michael@mgt01: ~  
michael@mgt01:~$ kubectl get node -owide  
NAME      STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE      KERNEL-VERSION   CONTAINER-RUNTIME  
virt1     Ready     <none>   4d1h  v1.29.7   172.16.1.60   <none>        Talos (v1.7.6) 6.6.43-talos    containerd://1.7.18  
virt2     Ready     <none>   4d1h  v1.29.7   172.16.1.61   <none>        Talos (v1.7.6) 6.6.43-talos    containerd://1.7.18  
virt3     Ready     <none>   4d1h  v1.29.7   172.16.1.62   <none>        Talos (v1.7.6) 6.6.43-talos    containerd://1.7.18  
virtcp    Ready     control-plane 4d3h  v1.29.7   172.16.1.59   <none>        Talos (v1.7.6) 6.6.43-talos    containerd://1.7.18  
michael@mgt01:~$
```



Homelab



What makes a good hypervisor?

- Hyper converged storage support
- VLAN / SDN support
- Live migration of virtual machines
- Templating
 - Golden images
 - cloud-init
- Snapshotting
- Memory sharing
- Overcommit on CPU and Memory
- A nice GUI
- API driven

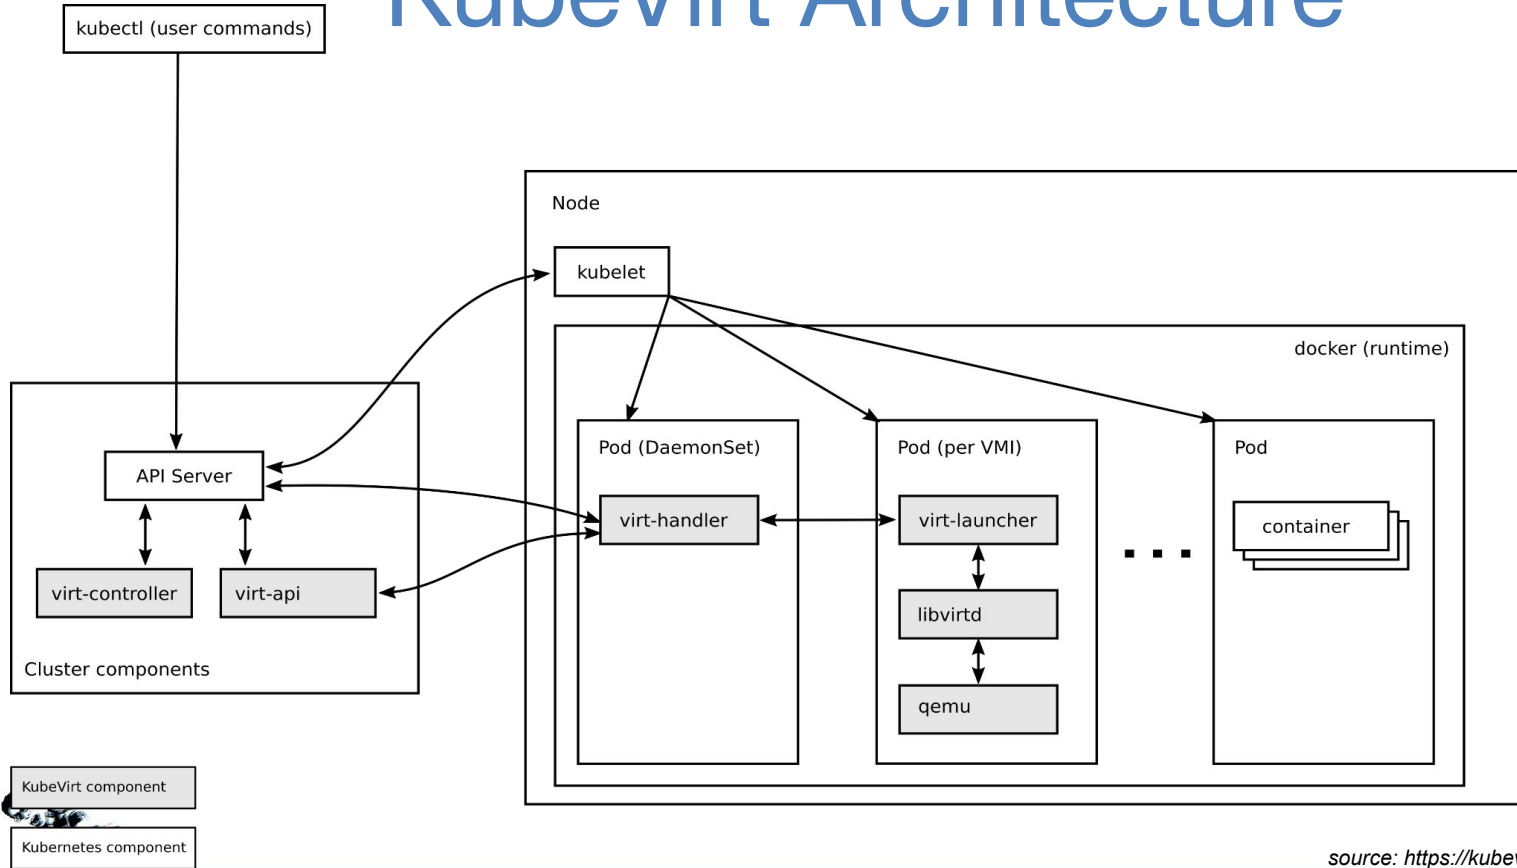


What is KubeVirt?

- Run virtual machines on k8s
- Extends k8s API with CRDs
- Uses Libvirt, qemu and kvm
- Run containers alongside virtual machines
- Core component for:
 - Harvester
 - Openshift virtualization
- Most commits in Github repo from Red Hat → 17487 (as of 5th of september)



KubeVirt Architecture



source: <https://kubevirt.io/user-guide/architecture/>

KubeVirt installation on Talos

Install Kubevirt operator:

```
# Point at latest release
$ export RELEASE=$(curl
https://storage.googleapis.com/kubevirt-prow/release/kubevirt/kubevirt/stable.txt)

# Deploy the KubeVirt operator

$ kubectl apply -f \
    https://github.com/kubevirt/kubevirt/releases/download/\${RELEASE}/kubevirt-operator.yaml
```



KubeVirt installation on Talos

apply Kubevirt CR:

```
---
apiVersion: kubevirt.io/v1
kind: KubeVirt
metadata:
  name: kubevirt
  namespace: kubevirt
spec:
  configuration:
    developerConfiguration:
      featureGates:
        - LiveMigration
        - NetworkBindingPlugins
        - Snapshot
    smbios:
      sku: "TalosCloud"
      version: "v0.1.0"
      manufacturer: "Talos Virtualization"
      product: "talosvm"
      family: "ccio"
```



KubeVirt installation on Talos

Install virtctl

```
# install virtctl with krew  
$ kubectl krew install virt
```

KubeVirt installation on Talos

Caveats:

- When using single disk nodes: make sure to upgrade with `talos upgrade --preserve=true`
- Make sure to set an exemption for the kubevirt namespace when using `PodSecurity`.
- When using Multus, make sure to configure your bridge properly:

```
network:  
  hostname: virt3.lan.alcatrash.net  
  interfaces:  
    - interface: br0  
      addresses:  
        - 172.16.1.62/24  
      bridge:  
        stp:  
          enabled: true  
        interfaces:  
          - eno1  
      routes:  
        - network: 0.0.0.0/0  
          gateway: 172.16.1.254
```



CDI: Containerized data importer

- CDI is used to import disks before the creation of a VM.
- Images supported are:
 - Qcow2
 - Raw
 - Iso
- Data sources where the images come from:
 - Upload from client
 - http/https
 - Container registry
 - Another pvc
- Create a CR called `DataVolume`
- `DataVolume` will create a PVC.



CDI: Containerized data importer

Install CDI operator:

```
# Point to latest release
$ export TAG=$(curl -s -w %{redirect_url} \
  https://github.com/kubevirt/containerized-data-importer/releases/latest)

$ export VERSION=$(echo ${TAG##*/})

# install operator
$ kubectl create -f \
  https://github.com/kubevirt/containerized-data-importer/releases/download/${VERSION}/cdi-operator.yaml
```


CDI: Containerized data importer

apply CDI CR:

```
---
apiVersion: cdi.kubevirt.io/v1beta1
kind: CDI
metadata:
  name: cdi
spec:
  config:
    scratchSpaceStorageClass: local-path
  featureGates:
    - HonorWaitForFirstConsumer
  podResourceRequirements:
    requests:
      cpu: "100m"
      memory: "60M"
    limits:
      cpu: "750m"
      memory: "2Gi"
```

CDI: Containerized data importer

Creating a DataVolume to import a base os disk:

```
apiVersion: cdi.kubevirt.io/v1beta1
kind: DataVolume
metadata:
  name: debian-12-image
  namespace: virtualmachines
spec:
  source:
    http:
      url:
"https://cloud.debian.org/images/cloud/bookworm/latest/debian-12-generic-amd64
.raw"
  pvc:
    accessModes:
      - ReadWriteMany
    resources:
      requests:
        storage: 3Gi
    storageClassName: nfs-client-zimaboard
```

CDI: Containerized data importer

Creating a Datavolume from imported disk:

```
apiVersion: cdi.kubevirt.io/v1beta1
kind: DataVolume
metadata:
  name: debian-external-pvc
  namespace: virtualmachines
spec:
  source:
    pvc:
      name: debian-12-image
      namespace: virtualmachines
  pvc:
    accessModes:
      - ReadWriteMany
    resources:
      requests:
        storage: 10Gi
    storageClassName: longhorn-rwx
```

CDI: Containerized data importer

Caveats

- Make sure to install the local path provisioner. If not installed, the importer pod will crash because it can't write any scratch space.
- Make sure to set proper limits in the CDI CR to allow the import to succeed. If not, your CDI importer pod will be `OOMKilled`.



But what about.... ?

- Vlan with and networking? → Multus
- Shared and Hyperconverged storage? → NFS and Longhorn
- Live migration? → Kubevirt Livemigration CR

Vlans and networking:

- Only use Multus when you want to expose your vm to the external network.
- First: make sure your bridge is configured properly
- Install Multus:
 - CNI plugin to attach multiple interfaces to pods
 - Make sure to patch the `DaemonSet`
- Install Whereabouts: Assigns ip addresses cluster wide to your VMs.
 - Only needed when using multiple nodes in your cluster.
- Create your `NetworkAttachmentDefinition` CR
- Connect your `VirtualMachine` NIC to that `NetworkAttachmentDefinition`



Vlans and networking:

Patch the Multus DaemonSet:

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: kube-multus-ds
  namespace: kube-system
spec:
  template:
    spec:
    ...
    volumes:
      - name: host-run-netns
        hostPath:
          path: /run/netns /var/run/netns
```



Vlans and networking:

Create a NetworkAttachmentDefinition:

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
  name: bridge-whereabouts
  namespace: virtualmachines
spec:
  config: '{
    "cniVersion": "0.3.1",
    "name": "bridge-whereabouts",
    "type": "bridge",
    "bridge": "br0",
    "ipam": {
      "type": "whereabouts",
      "range": "172.16.1.0/24",
      "range_start": "172.16.1.120",
      "range_end": "172.16.1.150",
      "gateway": "172.16.1.254",
      "routes": [
        { "dst": "0.0.0.0/0" }
      ]
    }
  }'
```


Vlans and networking:

Attaching the NetworkAttachmentDefinition to your VM:

```
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: debian-external-vm
spec:
  template:
    spec:
      domain:
        ...
        interfaces:
          - name: external
            bridge: {}
      networks:
        - name: external
      multus:
        default: true
        networkName: virtualmachines/bridge-whereabouts
```



Shared and hyper converged storage:

- Local path provisioner → Only used for temporary storage for CDI
- NFS CSI or NFS Subdir provisioner → Only for ISO's and disk images
- Longhorn for Hyperconverged storage
- Why not use Rook Ceph?
 - Single disk is not supported
- For Longhorn: Create a storage class with **ReadWriteMany** → used for Live migration



Shared and hyper converged storage:

Create Longhorn StorageClass with ReadWriteMany:

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: longhorn-rwx
provisioner: driver.longhorn.io
allowVolumeExpansion: true
reclaimPolicy: Delete
volumeBindingMode: Immediate
parameters:
  numberOfReplicas: "3"
  staleReplicaTimeout: "2880"
  fromBackup: ""
  fsType: "ext4"
  nfsOptions: "vers=4.2,noresvport,softerr,timeo=600,retrans=5"
```



Shared and hyper converged storage:

Some more information:

- Longhorn uses ISCSI for `ReadWriteOnce`
 - Make sure to install iscsi extension
- Longhorn creates a NFS Server per replica item when using `ReadWriteMany` storageclass.
- NFS Server based on the Ganesha project
- CSI provisioner will create a PVC on that NFS Server
- Make sure to configure a bind mount to `/var/lib/longhorn` on your worker nodes.



Live Migration

- Only available when:
 - Using PVC's with `ReadWriteMany` `StorageClass`
 - `LiveMigration` is enabled in the `FeatureGate` of the `Kubevirt` CR
- Can be initiated when executing: **`kubectl virt migrate <virtualmachinename>`**
- Creates a new CR called `VirtualMachineInstanceMigration`
- When using different cpu's in your cluster, make sure to set a CPU type in your VM
 - Kubevirt sets labels with `cpu-model-migration.node.kubevirt.io/<cpu-type>`. Choose the CPU type based on those labels.



Live Migration

Set CPU type:

```
---
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
  name: debian-external-vm
spec:
  running: true
  template:
    spec:
      domain:
        cpu:
          cores: 2
          model: Haswell-noTSX-IBRS
```



What about... ?

- Templating → Supported with the `VirtualMachineClone` CR
- Snapshotting → Supported with the `VirtualMachineSnapshot` CR
 - CSI must support `VolumeSnapshotClasses`
- Memory sharing and overcommitment:
 - KSM (Memory sharing) not supported on Talos
 - `CONFIG_KSM` not enabled in kernel
 - Overcommitment is supported in Kubevirt → beta.



Demo time

- Creating a virtual machine
- Migrating a virtual machine to another node

Demo - Create a Virtual machine

```
fedora-vm-pvc      CloneInProgress 69.86%      82s
fedora-vm-pvc      CloneInProgress 70.36%      84s
fedora-vm-pvc      CloneInProgress 76.76%      86s
fedora-vm-pvc      CloneInProgress 74.82%      88s
fedora-vm-pvc      CloneInProgress 79.26%      90s
fedora-vm-pvc      CloneInProgress 79.86%      92s
fedora-vm-pvc      CloneInProgress 80.20%      94s
fedora-vm-pvc      CloneInProgress 84.27%      96s
fedora-vm-pvc      CloneInProgress 88.72%      98s
fedora-vm-pvc      CloneInProgress 93.18%      100s
fedora-vm-pvc      CloneInProgress 97.57%      102s
fedora-vm-pvc      CloneInProgress 100.00%      104s
fedora-vm-pvc      CloneInProgress 100.00%      106s
fedora-vm-pvc      CloneInProgress 100.00%      108s
fedora-vm-pvc      CloneInProgress 100.00%      110s
fedora-vm-pvc      CloneInProgress 100.00%      112s
fedora-vm-pvc      Succeeded      100.0%      113s
^C$ kubectl virt start fedora-vm
VM fedora-vm was scheduled to start
$ kubectl get vel -o wide -w
NAME      AGE    PHASE      IP            NODENAME      READY    LIVE-MIGRATABLE  PAUSED
fedora-vm 6s     Scheduled  virti         virti         False
fedora-vm 6s     Scheduled  virti         virti         False    True
fedora-vm 6s     Running   10.244.8.233  virti         False   False    True
fedora-vm 6s     Running   10.244.8.233  virti         True    True     True
fedora-vm 6s     Running   10.244.8.233  virti         True    True     True
```



```

$ kubectl get pods
NAME                                READY     STATUS    RESTARTS   AGE
fedora-vm                           1/1       Running   0           1m
fedora-vm                           1/1       Running   0           1m
fedora-vm                           1/1       Running   0           1m
fedora-vm                           1/1       Running   0           1m
fedora-vm                           1/1       Running   0           1m
fedora-vm                           1/1       Running   0           1m

$ kubectl cp /dev/null /dev/null --dry-run
Warning: Pod 'fedora-vm' is not in a state that allows migration.

```



Takeaways

- Talos is awesome!
- KubeVirt is awesome!
- No enterprise grade GUI available:
 - KubeVirt Manager is the most promising
 - You could use Openshift Console → Does work for basic virtual machine mgmt
- Steep learning curve
- A lot of moving parts
- Watch out for caveats:
 - Multus → patch the `DaemonSet`
 - CDI → Local path storage to write scratch space
 - Configure your bridge properly in Talos
 - Upgrade with `--preserve=true`. If not, Longhorn storage is gone! :(



Conclusion

Enterprise grade production ready? **NO:**

- No KSM Support at this moment
- Overcommitment still in beta phase
- No enterprise grade GUI available

Startup grade production ready? **YES:**

- When staff is qualified enough to use Kubernetes / Talos / Kubevirt
- When not relying on GUI and overcommitment of CPU and memory
- Calculating a steep learning curve when designing a Kubevirt cluster



Questions?

Where to find me:

- Blog: <https://michaeltrip.nl>
- Email: m.trip@ATComputing.nl or michael@alcatrash.org
- Github: <http://github.com/michaeltrip>
- LinkedIn: <https://www.linkedin.com/in/michaeltrip/>



<https://github.com/michaeltrip/taloscon2024/>